

# Deep Learning Assignment II - Report

Chen-Yang Yu,

AI Robotics, National Cheng Kung University

Jun 20, 2024

<https://github.com/LittleFish-Coder/deep-learning-assignment-2>

## 1. Dataset:

Mini-ImageNet is a subset of the ImageNet dataset, which contains 50 classes.

- Training Set: 63325
- Validation Set: 450
- Testing Set: 450

Sample images from the dataset:

n02111277_1207.JPEG	n02112137_7.JPEG	n02112137_112.JPEG
		

## Data Preprocessing

It's worth noting that each image does not have the same size, so we need to standardize the size of all images. Moreover, some images are in RGB format, while others are in grayscale format.

- Image Size: we resize all images to **256x256** pixels.
- Image Channel: we convert all images to **RGB format (3 channels)**.

## Evaluation Performance

During training, we will evaluate the model using the validation set. Performance metrics such as test accuracy and loss will be discussed in the next 2 sections.

<https://github.com/LittleFish-Coder/deep-learning-assignment-2>

## 2. Task 1: Designing a Convolution Module for Variable Input Channels

We will discuss the 2 models' performance for dynamic input channels image.

### 2.1. Hyperparameters

- Epoch: 25
- Learning Rate: 0.001
- Batch Size: 64
- Input Size: (3, 256, 256)
- Optimizer: Adam
- Criterion: CrossEntropyLoss

### 2.2. Simple CNN

- Model Architecture:

Layer (type)	Output Shape	Param #
Conv2d-1	[-1, 32, 256, 256]	896
BatchNorm2d-2	[-1, 32, 256, 256]	64
MaxPool2d-3	[-1, 32, 128, 128]	0
Conv2d-4	[-1, 64, 128, 128]	18,496
BatchNorm2d-5	[-1, 64, 128, 128]	128
MaxPool2d-6	[-1, 64, 64, 64]	0
Conv2d-7	[-1, 128, 64, 64]	73,856
BatchNorm2d-8	[-1, 128, 64, 64]	256
MaxPool2d-9	[-1, 128, 32, 32]	0
AdaptiveAvgPool2d-10	[-1, 128, 1, 1]	0
Linear-11	[-1, 256]	33,024
Dropout-12	[-1, 256]	0
Linear-13	[-1, 50]	12,850

Total params: 139,570

Trainable params: 139,570

Non-trainable params: 0

-----  
Input size (MB): 0.75

Forward/backward pass size (MB): 63.01


Params size (MB): 0.53

Estimated Total Size (MB): 64.29  
-----

- **Validation Accuracy & Loss**

We save the best model based on the validation loss

- Best Validation Accuracy: 0.3379
- Best Validation Loss: 0.0349



- **Test Accuracy & Loss**

	RGB	RG	GB	R	G	B
Accuracy	0.3556	0.3111	0.0378	0.0222	0.0222	0.02
Loss	0.0412	0.0479	0.1202	0.2145	0.1152	0.2187

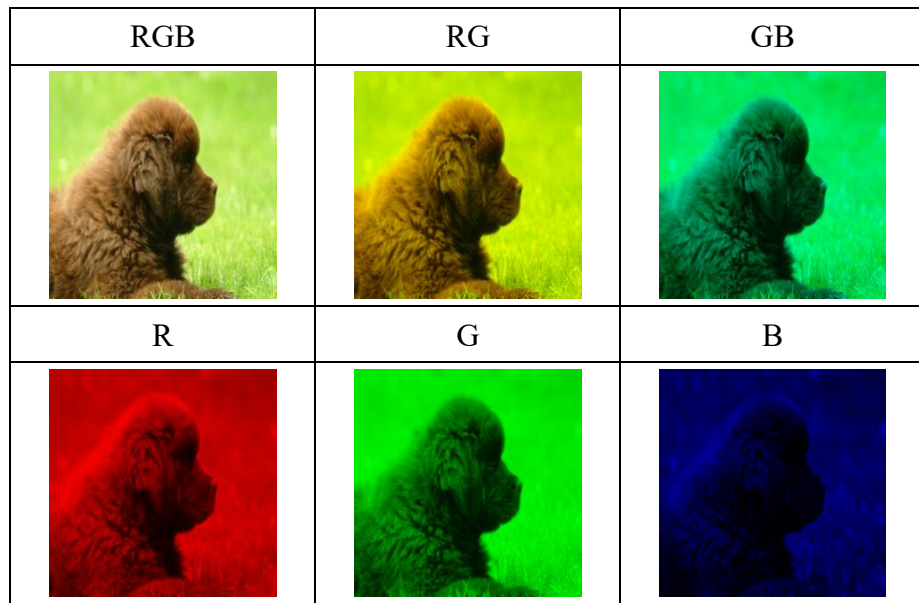
- **Dataset Design**

We will test the model on 6 downstream tasks, each using different combinations of the RGB channels.

Since the simple CNN can only process images with 3 channels, we have redesigned the test dataset to accommodate this limitation. When testing different channel combinations, we will first select the desired channels and then copy these channels to a new image. The channels that are not selected will be padded with zeros. This ensures that the image remains in a 3-channel format.

For example, if the 'RG' combination is selected, we will copy the 'R' and 'G' channels from the raw image. The 'B' channel will be set to zero.

Below are samples for images with each selected channel combination:



### 2.3. Dynamic CNN

#### ● Model Architecture

Layer (type)	Output Shape	Param #
Conv2d-1	[-1, 16, 254, 254]	160
ReLU-2	[-1, 16, 254, 254]	0
Conv2d-3	[-1, 1, 252, 252]	145
Conv2d-4	[-1, 16, 254, 254]	160
ReLU-5	[-1, 16, 254, 254]	0
Conv2d-6	[-1, 1, 252, 252]	145
PoolChannelAttention-7	[-1, 1, 252, 252]	0
Conv2d-8	[-1, 32, 252, 252]	320
BatchNorm2d-9	[-1, 32, 252, 252]	64
MaxPool2d-10	[-1, 32, 126, 126]	0
Conv2d-11	[-1, 64, 126, 126]	18,496
BatchNorm2d-12	[-1, 64, 126, 126]	128
MaxPool2d-13	[-1, 64, 63, 63]	0
Conv2d-14	[-1, 128, 63, 63]	73,856
BatchNorm2d-15	[-1, 128, 63, 63]	256
MaxPool2d-16	[-1, 128, 31, 31]	0
AdaptiveAvgPool2d-17	[-1, 128, 1, 1]	0
Linear-18	[-1, 256]	33,024
Dropout-19	[-1, 256]	0
Linear-20	[-1, 50]	12,850
CNN-21	[-1, 50]	0

Total params: 139,604

Trainable params: 139,604

Non-trainable params: 0

<https://github.com/LittleFish-Coder/deep-learning-assignment-2>

---


Input size (MB): 0.75  
Forward/backward pass size (MB): 93.98  
Params size (MB): 0.53  
Estimated Total Size (MB): 95.26

---

- **Validation Accuracy & Loss**

We save the best model based on the validation loss

- Best Validation Accuracy: 0.2889
- Best Validation Loss: 0.0410



- **Test Accuracy & Loss**

	RGB	RG	GB	R	G	B
Accuracy	0.3044	0.3111	0.2911	0.3089	0.3089	0.2467
Loss	0.0480	0.0479	0.0483	0.0489	0.0478	0.0526

- **Dataset Design**

The Dynamic CNN can handle different input channels. Therefore, if 'RG' is selected for testing, we will create a new image containing only the 'RG' channel information from the raw image.

In this way, we can create images with the following dimensions:

- (1, height, width) for 'R', 'G', or 'B'
- (2, height, width) for 'RG' or 'GB'
- (3, height, width) for 'RGB'

## 2.4. Comparison

By comparing the results from the testing phase of these two models, the dynamic CNN outperforms the simple one on the 6 downstream tasks.

## 3. Task 2: Designing a Two-Layer Network for Image Classification

### 3.1. Hyperparameters

- Epoch: 50
- Learning Rate: 0.001
- Batch Size: 64
- Input Size: (3, 256, 256)
- Optimizer: Adam
- Criterion: CrossEntropyLoss


### 3.2. ResNet34

- **Model Architecture**

-----  
Total params: 21,310,322  
Trainable params: 21,310,322  
Non-trainable params: 0  
-----

Input size (MB): 0.75  
Forward/backward pass size (MB): 125.75  
Params size (MB): 81.29  
Estimated Total Size (MB): 207.80  
-----

- **Validation Accuracy & Loss:**
  - Best Validation Accuracy: 0.5622
  - Best Validation Loss: 0.0212



- **Test Accuracy & Loss**

- Test Accuracy: 0.5978
- Test Loss: 0.0221

### 3.3. Attention CNN

- **Model Architecture**

Layer (type)	Output Shape	Param #
Conv2d-1	[-1, 32, 256, 256]	896
BatchNorm2d-2	[-1, 32, 256, 256]	64
ReLU-3	[-1, 32, 256, 256]	0
Conv2d-4	[-1, 32, 256, 256]	128
BatchNorm2d-5	[-1, 32, 256, 256]	64
ReLU-6	[-1, 32, 256, 256]	0
Block-7	[-1, 32, 256, 256]	0
MaxPool2d-8	[-1, 32, 128, 128]	0
Conv2d-9	[-1, 64, 128, 128]	18,496
BatchNorm2d-10	[-1, 64, 128, 128]	128
ReLU-11	[-1, 64, 128, 128]	0
Conv2d-12	[-1, 64, 128, 128]	2,112
BatchNorm2d-13	[-1, 64, 128, 128]	128
ReLU-14	[-1, 64, 128, 128]	0
Block-15	[-1, 64, 128, 128]	0
MaxPool2d-16	[-1, 64, 64, 64]	0
Conv2d-17	[-1, 128, 64, 64]	73,856
BatchNorm2d-18	[-1, 128, 64, 64]	256
ReLU-19	[-1, 128, 64, 64]	0
Conv2d-20	[-1, 128, 64, 64]	8,320
BatchNorm2d-21	[-1, 128, 64, 64]	256
ReLU-22	[-1, 128, 64, 64]	0
Block-23	[-1, 128, 64, 64]	0
MaxPool2d-24	[-1, 128, 32, 32]	0
Conv2d-25	[-1, 256, 32, 32]	295,168
BatchNorm2d-26	[-1, 256, 32, 32]	512
ReLU-27	[-1, 256, 32, 32]	0
Conv2d-28	[-1, 256, 32, 32]	33,024
BatchNorm2d-29	[-1, 256, 32, 32]	512
ReLU-30	[-1, 256, 32, 32]	0
Block-31	[-1, 256, 32, 32]	0
MaxPool2d-32	[-1, 256, 16, 16]	0
Conv2d-33	[-1, 1, 16, 16]	98
Sigmoid-34	[-1, 1, 16, 16]	0
SpatialAttention-35	[-1, 1, 16, 16]	0
Dropout-36	[-1, 256, 16, 16]	0
AdaptiveAvgPool2d-37	[-1, 256, 1, 1]	0
Linear-38	[-1, 50]	12,850


Total params: 446,868  
Trainable params: 446,868  
Non-trainable params: 0

---

Input size (MB): 0.75  
Forward/backward pass size (MB): 218.01  
Params size (MB): 1.70  
Estimated Total Size (MB): 220.46

- **Validation Accuracy & Loss**

- Best Validation Accuracy: 0.5222
- Best Validation Loss: 0.0235



- **Test Accuracy & Loss**

- Test Accuracy: 0.5911
- Test Loss: 0.0239

### 3.4. Comparison

By adding a Spatial Attention module to a simple CNN with some residual blocks, we are able to rival the ResNet34 model.

	ResNet34	Attention CNN
Test Accuracy	0.5978	0.5911
Test Loss	0.0221	0.0239
Total params	21,310,322	446,868